


A N A T L A S O F T H E A G E N T S T R A T A

THE AGENT, SEEN WHOLE

The Plates

Twelve plates in the order they should be experienced —
synthesis first, then the agent column, then the frameworks.



Compiled for A. Yedi · Cycle MMXXVI · Revision I

2026-05-14

S U B S T R A T E · V O L U M E I I I · M A S T E R P L A T E S

Twelve plates in the order they should be read

S I - T H E S Y N T H E S I S

M Master Plate · The agent layer as a single connected system

Begin here — every later plate is a zoom into one quadrant of this one.

S I I - T H E A G E N T C O L U M N · S U B S T R A T E V O L I

I Index of the agent column — the ten sub-strata + four meta

II The cognition loop — perceive · plan · act · reflect

III Memory and state — short-term, episodic, long-term, shared

IV Action surfaces — tool use, web, code, computer, voice

V Vertical products — where the agent meets the buyer

VI Evaluation, observability, safety — the meta-strata

S I I I - T H E F R A M E W O R K P L A T E S · S U B S T R A T E V O L I I

VII OCQ heat map — opportunities, challenges, open questions

VIII Wardley map — evolution of every component, claimable flags

IX Helmer's 7 Powers — which powers compound at the agent layer

X JTBD — the six ecosystem-level jobs the agent layer performs

XI Action Portfolio — the seven Bets refreshed, sequenced, sized

T H E S Y N T H E S I S F R O N T I S P I E C E

Master Plate

Begin here · the agent seen whole

§ I of III

S U B S T R A T E · V O L U M E I I I · M A S T E R P L A T E S

Stratum agenticum tota machina · quomodo strata simul operentur



T H E A G E N T C O L U M N

Plates I–VI

Sub-strata + the meta · the agent layer unfolded

§ II of III

S U B S T R A T E · V O L U M E I I I · M A S T E R P L A T E S

THE AGENT COLUMN

A higher-resolution zoom into Stratum XIII of the parent atlas — the agent layer unfolded as ten sub-strata plus four wrapping meta-strata.
Stratum agenticum — denso compositum, ordo descendens, gradibus decem cum quattuor strata circumdantia

X SUB-STRATA · IV META · DESCENDING

X	END-USER · SURFACE form factors	chat · cli · ide · trigger · slack · whatsapp · voice · ar	10
I X	VERTICAL · PRODUCTS productized agents	sierra · decagon · glean · harvey · hippocratic · clay · rogo	9
V I I I	RUNTIME · SAFETY tool-boundary defense	lakera · nemo · llama-guard · bedrock · purview	8
V I I	EVAL · OBSERVABILITY trajectory analytics	braintrust · langsmith · langfuse · arize · galileo · inspect	7
V I	ACTION · SURFACES agent reaches into world	e2b · modal · browserbase · computer-use · vapi	6
V	PLANNING · REASONING cognition loop	react · planner-executor · o-series · extended-thinking	5
I V	MEMORY · STATE persistence	mem0 · letta · zep · graphrag · lab-native	4
I I I	TOOL · PROTOCOL interface to world	mcp · function calling · gateways · a2a	3
I I	AGENT · RUNTIMES harness	claude agent sdk · openai agents · adk · langgraph · mastra	2
I	FOUNDATION · CAPABILITY the agentic delta	claude · gpt-5 · gemini · deepseek · tau · swe · osworld	1

M E T A · W R A P P I N G F O R C E S

A	CAPABILITY · SAFETY	asl · preparedness · frontier · aisi · metr
B	REGULATION	eu ai act · ca sb 53 · sectoral · bartz · nyt
C	ECONOMICS	per-task · outcome pricing · finops
D	GEOPOLITICS	sovereign agents · export controls · jurisdiction

K E Y	sub-stratum	component	meta-stratum	dependency
-------	-------------	-----------	--------------	------------

CAPABILITY · SUBSTRATE

Strata I–III · the model, the runtime, the protocol — the geological basement on which every agent runs.
Stratum primum, secundum, tertium · capacitas, harnesia, protocollum tooluse

III

TOOL · PROTOCOL

MODEL CONTEXT PROTOCOL · ADOPTION · FORK VECTORS

● MCP — Anthropic Nov '24 JSON-RPC; Py + TS SDKs	● OpenAI adoption — Mar '25 Agents SDK · Responses API · mini · Vertex AI · A2A Copilot Studio · SK	● Google adoption — Apr '25	● Microsoft adoption — May '25
● Hugging Face — Jun '25 MCP-compatible Spaces	● Linux Foundation — Dec '25 Agentic AI Foundation	● Registry · ~11.4K servers Apr '26 LF stats	● Cloudflare Workers AI MCP Dec '24 · gateway
● Kong MCP Gateway GA Jul '25 · F500 logos	● Pomerium identity proxy May '25	● Function calling (legacy) OpenAI Jun '23	● Structured outputs OpenAI Aug '24
● Stripe MCP first-party · canonical	● Linear · Cloudflare MCP first-party tier	● GitHub MCP MSFT canonical	● A2A protocol — Apr '25 agent-to-agent (Google)
● Fork vector — proprietary tool calls OAI latency wins	● Fork vector — quality split long tail abandonware		

stratum tertium · the joint to the world

MCP · JSON-RPC · GATEWAYS

REGISTRY · MAY '26

~11.4K

HYPERSCALERS COMMITTED

4 / 4

GATEWAY CATEGORY

GA

II

RUNTIMES

AGENT RUNTIMES · HARNESESSES · ORCHESTRATORS

● Claude Agent SDK Sep '25 · sub-agents · small	● OpenAI Agents SDK Mar '25 · handoffs	● Google ADK Apr '25 · Vertex Agent Ecosystem	● LangGraph 1.0 May '25 · production
● LangSmith obs layer · margin lives	● Mastra TypeScript · WC24 · Vertex	● Pydantic AI Pydantic validation · Logfire	● CrewAI multi-agent · \$18M Oct '24
● Smolagents (HF) minimalist · Dec '24	● AutoGen / AG2 MSR + community fork	● DSPy · BAML compile-time / IR layer	● Inspect AI (UK AISI) eval framework · OSS
● Mistral Agents API May '25 · EU	● Manus AI (CN) Butterfly Effect · viral Q12526	● Sarvam AI (IN)	● Aleph Alpha compliance-first · DE

stratum secundum · the harness

SDK · LOOP · SUB-AGENTS

SHIPPED · LAST 14 MO

4 SDKs

LANGGRAPH ARR

\$40–60M (inf)

RUNTIME GROSS MARGIN

THINNEST

I

FOUNDATION · CAPABILITY

AGENTIC CAPABILITY DELTAS · NOT PROVIDER ECONOMICS

● Claude Opus 4.5 Nov '25 · 82.8% SWE-V	● Claude Sonnet 4.5 interactive executor	● GPT-5 / 5.1 thinking ~80% SWE-V w/ tools	● ChatGPT Agent Operator+DR · OSWorld classm
● Gemini 2.5 Pro Deep Think GAIA ~70 · BrowseComp ~50	● Project Mariner Apr '26 · OSWorld ~38	● DeepSeek R1 / R2 open weights · mid-50s SWE-V	● Qwen QwQ / Qwen3 open · CN
● Kimi K1.5 / K2 long-context CN · Moonshot	● Meta Llama 4 Apr '25 · -10–15 pts vs	● Anthropic sub-agents children · 6mo context degr. past 32K tokens	● Reasoning regression
● Error compounding ceiling 0.95*20 = 36%	● Computer-use ceiling OSWorld ~50% vs human	● Extended thinking 64K 72% interactive vs back-official	● Tool use in thinking GPT-5 · Deep Think

stratum primum · the agentic delta

TAU · SWE · OSWORLD · BROWSECOMP

SWE-BENCH VERIFIED

82.8%

OSWORLD (MAY '26)

~50%

REASONING-MODEL TIERS

9

COGNITION · LOOPS

Strata IV–V · the inner cognition of the working agent — what it remembers and how it plans the next step.
Stratum quartum et quintum · memoria et ratio · circuitus interni

V

PLANNING · REASONING

WHAT SURVIVED · TEST-TIME COMPUTE · ECONOMICS

● ReAct (Yao Oct '22) still dominant 2026 loop prompting · alive	● Plan-and-Solve (May '23) lab curiosity · rare prov verifier sub-agent · alive	● Tree-of-Thoughts (May '23) emerging 2025–26	● Reflexion (Mar '23) reasoning model tier opens
● AutoGPT / BabyAGI dead in prod · idea survives	● Planner-executor split adaptable architectural id	● Meta-planning / hierarchical	● o1 / o3 (Sep–Dec '24)
● Claude extended thinking Feb '25 · 64K budgets	● Gemini Deep Think (May '25) super-linear on hard	● DeepSeek R1 / R1-Zero Jan '25 · pure RLVR	● Qwen QwQ-32B · Kimi K1.5 open reasoning
● Tool use IN thinking outer ReAct collapses	● `reasoning_effort` knob OpenAI · buyer-controlled	● PRM (Lightman May '23) n-model · not exposed	● GRPO at trajectory DeepSeek · default RL

stratum quintum · the inner loop

REACT · PLANNER-EXECUTOR · TTC

CX TRAJECTORY · TOKENS

4–25K

CODE AGENT · TOKENS

30–250K

DEEP RESEARCH · TOKENS

0.5–5M

IV

MEMORY · STATE

STANDALONE OR ABSORBED · COMPLIANCE WEDGE

● Memo0 (YC W24) \$24M A Dec '25 · 40K state \$10M seed · A rumored '26	● Letta (MemGPT) \$15M A Feb '25 · HIPAA	● Zep — Graphiti OSS pipeline · seed Dec '24	● Cognee Sep '25 paper · CN
● Microsoft GraphRAG Apr '24 paper · Jul '24 @ \$1M '24 · cheaper graph	● LightRAG (HKU) research-only	● HippoRAG · PathRAG · OG-RAG	● MemoryOS (Tsinghua)
● ChatGPT memory Apr '24 paid · Sep '24 free	● Claude Projects memory Apr '26	● Gemini personalization GA Mar '26	● Working memory (LangGraph) engineering not category
● Procurement — Article 17 Mem0 forget API Feb '26 Athelas · F500 health	● HIPAA story (Zep)	● Contextual integrity Nissenbaum · no vendor	● Falsifiability test 70% pitches = RAG + writes

stratum quartum · persistence

EPISODIC · SEMANTIC · PROCEDURAL

STANDALONE CEILING (INF) LAB MEMORY MAU

\$200–400M

~800M

CRUX #5 (INHERITED)

OPEN

ACTION · DEFENSE

Strata VI–VIII · where the agent reaches into the world, how we observe that it did the right thing, and what defends when it does not.
Stratum sextum, septimum, octavum · operatio, observatio, custodia

VIII

RUNTIME · SAFETY

GUARDRAILS · INJECTION DEFENSE · BUNDLING

- Lakera (Guard · Red)
Series A '23 · B Q4 '25
 - Apex Security
Sequoia seed Mar '24
 - Purview AI Hub · Defender
MSFT Ignite Nov '24
 - Promptfoo (now OAI)
red-team OSS · Sep '25
- Protect AI · to Palo Alto
~\$700M Q3 '25 · the bell Aug '24 · AI Defense
 - CalypsoAI
\$23M B · federal angle
 - Vertex AI Model Armor
Cloud Next '24
 - Garak (NVIDIA OSS)
vulnerability scanner
- Robust Intelligence · to Cisco
KPMG spin · GRC-leaning Apr '24 · Auto Reasoning
 - Cranium
GA Apr '26 · OSS
 - NeMo Guardrails 1.0
EchoLeak CVE class Mar '26 · indirect PI
- HiddenLayer
\$50M B Mar '24
 - AWS Bedrock Guardrails
Apr '24 · Auto Reasoning
 - Llama Guard 3 · Prompt Guard 2
Meta open weights
 - Constitutional classifiers
Anthropic · baked in

stratum octavum · the tool boundary

PI · ACTION GATES · SANDBOXING

PURE · PLAYS ABSORBED

18 MO

INDIRECT PI

UNSOLVED

DEFENSE STACK · LAYERS

5

VII

EVAL · OBSERVABILITY

AGENT EVAL · OBS · BENCHMARKS

- Braintrust
\$36M A · B rumored '26
 - Galileo
\$45M B · C reported Apr U26 · FSI
 - AgentOps
CrewAI / AutoGen ecosystem
assessing as native matures
 - METR · time-horizon
Mar '25 · cite not procu
- LangSmith / LangChain
\$25M C · LangGraph 1.0
 - HumanLoop
U26 · FSI
 - Helicone · proxy
\$derra-co · CX proxy
 - TAU-bench / tau-2
- Langfuse
OSS · YC W23 · EU choice
 - Patronus AI · Lynx · Glider
\$17M A · defensible judg
 - Promptfoo (acq OAI)
Sep '25 · neutrality gon
 - SWE-Bench Verified · Live
code agency · ~75-80%
- Arize Phoenix / AX
\$70M C Dec '24
 - Comet Opik · DeepEval
@SS · pytest-style
 - Inspect AI (UK AISI)
EU AI Act conformity
 - OSWorld · WebArena · GAIA
computer/web/research

stratum septimum · trajectory analytics

TRAJECTORY · REPLAY · ONLINE

OTEL GENAI CONVENTIONS

JAN '26 GA

CONSOLIDATION (2027)

PARTIAL

PROCUREMENT GAP

OPEN

VI

ACTION · SURFACES

SANDBOXED CODE · BROWSER · COMPUTER · VOICE OUTBOUND

- E2B
OSS-core · ~\$15-25M ARR
 - Replit Agent sandbox
bundled · prosumer
 - Playwright + AI
the boring baseline
 - Google Project Mariner
GA Apr '26 · Gemini Ent
 - ElevenLabs Conv. AI
voice up-stack
- Modal
agent-workload pivot ·
 - HF Spaces · Codespaces
ad-hoc · weaker fit
 - Anchor · Hyperbrowser
MCP-compat 2025
 - MSFT Copilot Vision
Windows-native
 - Cartesia Sonic-2
~40ms first-byte
- Daytona
\$50M R · price-aggressive
 - Browserbase + Stagehand
>\$50M ARR Q1 '26
 - Anthropic Computer Use
Oct '24 · Claude for Chr
 - Vapi · Retell · Bland
voice trio · \$20M+ ARR
 - Deepgram Nova-3
ASR · accent
- Vercel Sandbox
GA Jan '26 · Firecracker
 - Browserless · Skyvern
incumbent · OSS
 - ChatGPT Agent
@p@rator rebrand mid '25
 - LiveKit Agents
eWebRTC substrate
 - Twilio · Telnyx · LiveKit Cloud
telephony

stratum sextum · agent reaches into world

SANDBOX · BROWSER · COMPUTER · VOICE

VOICE · TO · VOICE FLOOR

500-800ms

SANDBOX TIMELINE

18-24 MO

CU PRODUCTION

NO

PRODUCTIZED · LAYER

Strata IX - X - where enterprise revenue lands and where the human meets the agent -- the productized application and its surface.

Stratum nonum et decimum - productum agenticum et facies hominis

X

END-USER · SURFACES

T W E L V E F O R M F A C T O R S · T H E B I N D I N G C O N S T R A I N T

● Chat — web/desktop ChatGPT 800M WAU · plateApple Intelligence slot Claude Code ~\$500M RR Cursor · Copilot · Windsurf	● Chat — mobile/on-device n8n · Zapier · cron · ArcClaude for Chrome · ComeCo\$Mlot 365 ~30M seats Glean wedge	● CLI /terminal Embedded in SaaS	● In-IDE inline Slack / Teams · headless
● Agentic trigger Meta API Jan '26 · 2.9B Advanced Voice · AirPodsRay-Ban >2M · Plaud >1M	● In-browser sidebar Voice — inbound	● Wearable · AR · ambient	● Slack / Teams · headless Glean wedge
● WhatsApp · SMS · email	● Voice — inbound	● Wearable · AR · ambient	● Computer use (UX)
● API · SDK · no-UI 2 / 3 OAI rev · majority Anthropic mobile gap	● WhatsApp Business API	● Coze on Feishu · WeChat	● Manus · Yellow.ai · Haptik
● Pew chat-churn 41%	● Anthropic mobile gap		
Jan '26 · 'chat solved' noS&I slot · no voice			

stratum decimum · form factor

B U Y E R · T H R E A T · A D O P T I O N

M I C R O S O F T C O P I L O T A R R

\$5B+

W H A T S A P P M A U

2.9B

B I N D I N G C O N S T R A I N T

FORM

IX

VERTICAL · PRODUCTS

E N T E R P R I S E R E V E N U E L A Y E R · B U Y I N G M O T I O N B Y D O M A I N

● Sierra — \$175M+ ARR \$10B (Mar '26) · NYC dua\$4.5B Sep '25 · NYC/SF	● Decagon — \$80M+ \$7.2B Sep '25 · NYC grow\$5@ Feb '26 · NYC office	● Glean — \$300M+ \$977M Nov '25 · enterpri\$@.9B Jul '25 · IDE lead\$5B Codeium absorbed Mar '26	● Harvey — \$100M+ \$3B Sep '25 · prosumer fastest EU SaaS · Stockh\$2m 75B Jun '25 · scribe NYC HQ · ex-2Sigma · raise
● Hippocratic — \$50M+ \$2B Feb '26 · NYC HQ	● Augment Code — \$40M \$977M Nov '25 · enterpri\$@.9B Jul '25 · IDE lead\$5B Codeium absorbed Mar '26	● Cursor — \$500M RR \$13B-NYC: 20B raise rumor	● Cognition (Devin) NYC Chelsea · creative
● Replit Agent — \$150M \$3B Sep '25 · prosumer fastest EU SaaS · Stockh\$2m 75B Jun '25 · scribe NYC HQ · ex-2Sigma · raise	● Lovable — \$80M Q1'26 fastest EU SaaS · Stockh\$2m 75B Jun '25 · scribe NYC HQ · ex-2Sigma · raise	● Abridge — \$120M+ \$13B-NYC: 20B raise rumor	● Hebbia — \$50M+ NYC Chelsea · creative
● Clay — \$80M+ NYC HQ · \$1.5B Jan '26	● Rogo — \$30M+ NYC HQ · \$400M Jan '26	● Ramp AI org \$13B-NYC: 20B raise rumor	● Runway — \$100M+ NYC Chelsea · creative
● ElevenLabs — \$200M+ voice platform · \$3.3B CX maturing/stalling	● Cresta · Ada · Intercom Fin CX maturing/stalling	● Parloa · PolyAI · Regal EU/UK + NYC sales voice legal long tail	● EvenUp · Eve · Spellbook legal long tail
● 11x · Artisan · AiSDR synthetic SDR · 11x flat6600M Mar '26 · EU sovereign	● Mistral Le Chat Ent. synthetic SDR · 11x flat6600M Mar '26 · EU sovereign		

stratum nonum · productized agents

C X · C O D E · L E G A L · H E A L T H · F I N A N C E

W I N N E R - T A K E - M O S T S E G M E N T S E T # 2 N Y C T A R G E T S

3 / 9

10+

O U T C O M E P R I C I N G (T A Y L O R)

CX ONLY

META · WRAPPING FORCES

Meta A–D · the four forces that wrap the agent layer: capability-safety regimes, regulation, economics, geopolitics.
Quattuor stratorum meta · circumscribunt agentcam stratam

D

GEOPOLITICS

WHERE AGENTS RUN · EXPORT CONTROLS · SIT-CONTROLS

- Mistral Le Chat Ent. €600M @ €11B Mar '26 Humain ties · sovereign AI 300/350 commits · \$10B \$41M · Q1 '26 raise rumor
- Krutrim (Ola IN) \$50M · political tailwind Butterfly · \$500M rumor ByteDance · parallel-track agent Q2 '26
- US Commerce auth — Nov '25 G42 + Humain chips
- Agent-capability export rule 2026–27 regulatory front AI Act + GDPR + Schrems
- Humain (KSA)
- Coze on Feishu
- EU residency lever
- Sarvam AI (IN)
- WeChat Mini-Programs
- Trajectory siting Where memory + actions reside

meta · sovereign agents

S O V E R E I G N · E X P O R T · J U R I S D I C T I O N

C

ECONOMICS

FROM PER-TOKEN TO PER-TASK · OUTCOME PRICING

- Token price · –4× / –10× Q1 '24 to Q1 '26
- CX trajectory · 4–25K tok ~0.02–0.30 / task
- Code agent · 30–250K tok ~0.20–3 / task
- Deep research · 0.5–5M tok ~1–20 / task
- Planner-executor split 70–85% on cheap tier
- Vercel AI Gateway routing first-class · GA '26
- AWS Bedrock · OpenRouter tier routing standard
- Outcome pricing — CX per-resolution · \$1–4 (Taylor)
- Outcome — healthcare per-encounter · per-call \$30/hr knowledge · financial –18 mo window
- FinOps for Tokens (Bet #4)
- Hyperscaler bundling Bedrock auto-optim threat
- Per-seat survivors

meta · the inversion

P E R - T A S K · O U T C O M E · F I N O P S

B

REGULATION

EU ARTICLE 14 · CA SB 53 · BARTZ · SECTORAL

- EU AI Act · Art. 55 GPAI enforceable Aug 2 '25
- EU AI Act · Art. 14 oversight draft Apr '26 · agent Q2 '26 · Inspect AI cited effective Jan 1 '26
- EU AI Act · conformity drafts
- CA SB 53 frontier disclosure
- NY SS641 · state-level advancing 2026
- Trump Dec '25 preemption EO courts through '26–'27
- Bartz v. Anthropic \$1.5B training-data settlement vendor terms repricing
- NYT v. OpenAI
- FINRA · OCC · BD rules finance agent fence
- HIPAA · FDA SaMD health agent fence
- TCPA · state consent voice agent fence
- EU AI Office structural regulator

meta · the gauntlet

E U A I A C T · C A S B 5 3 · S E C T O R A L

A

CAPABILITY · SAFETY

LAB-SIDE SAFETY REGIMES · EVAL INSTITUTES

- Anthropic ASL framework ASL-3 Opus 4 May '25
- Anthropic RSP — Nov '24 binding policy
- OpenAI Preparedness Framework tracked categories
- Preparedness adjustment clause Risk #3 · open question
- Google Frontier Safety May '24 · maturing
- DeepMind capability evals auto-replication · cyber Nov '23 · Inspect framework early '26 · narrower scope
- UK AISI
- US AISI (post-restruct)
- METR (research org) time-horizon · 7-mo double-blind regulatory counterpart
- EU AI Office 2026–27 question
- ASL-4 capability emergence
- Model cards for agents in-flight standardization

meta · the deploy gate

A S L · P R E P A R E D N E S S · F R O N T I E R

T H E F R A M E W O R K P L A T E S


Plates VII–XI

OCQ · Wardley · 7 Powers · JTBD · Action Portfolio

§ III of III

S U B S T R A T E · V O L U M E I I I · M A S T E R P L A T E S


Fourteen agent strata × three lenses · scored opportunity, challenge, open question.
Stratum agenticum × instrumentum analyticum · conviction across the column



Four anchored agent user-needs dependencies cascading · evolution stage along the X-axis.
Cartographia agenticum · stage of evolution in 2025–2026




Per-stratum mapping of the seven Helmer powers · who holds them · direction of travel.
Septem potestates Helmeri applicatae stratis agenticiis



Seven ecosystem-level jobs the agent stack is hired for · Ulwick eight-phase map · Conclude is universally unserved.

Septem opera quaesita · phases octo · Concludere universaliter desertum



Seven refreshed Bets sequenced across three windows. Bet #1 first (claim power) · Bet #2 second (collect equity) · Bet #3 third (compound).
Septem electa per tres fenestras temporis · ordo agendi

B E T

MONTHS 0 - 6 · P L A N T

MONTHS 6 - 12 · C O M M I T

MONTHS 12 - 18 · C O M P O U N D

1

Procurement Operating Standard

5/5

Publish open Playbook (wk 12)

30 expert interviews

AI Gov sign-off audit

3-5 paid engagements

Procurement Audit reports

Article 14 playbook

SaaS productization decision

\$250-500K annual rev

Three-product bundle

2

Vertical Agent GTM Role

5/5

NYC search - 10 targets

Time offer before up-round

Consulting alumni pipeline

Sign at chosen vertical

Close Confirm + Conclude gap

Anthropic ARR resolution

Deliver job - vertical

Compound from inside

Re-rank by Crux #1 outcome

3

MCP-Native Practice

4/5

Audit 10 SaaS - gateway pair

Pair w/ Cloudflare - Kong

Reframe to advisory

Productize or sustain advisory

Gateway-adjacent positioning

Watch fork Crux #2

Compound - advisory + Playbook

Decide on platform vs. practice

4

FinOps for Trajectories

4/5

Free first audit - 5 cos

Case study

Per-trajectory rubric

Fold into Bet #1 module

Decide scale or absorb

Decay window honest

Bedrock auto-routing watch

5

RAG + Memory Architecture

4/5

Bundle as audit

Memory architecture line

Vendor-neutral guidance

3 architecture audits

Three-product practice

Sustain - selective

Track CRUX #5

6

Operator Newsletter

3/5

Kit v1 - wk 4

Cadence decision - wk 8

JTBD Job 5 angle

5K subs / 3 inbound - or kill

Conclude-phase translation

Sustain - paid tier

Distribution layer for #1-#3

7

VC Operating Partner

3/5

Background networking

RAAIS - Betaworks - MAD

No applications

3 NYC funds - relationships

Watch principal-to-operator

Primary path if 1-3 stall

Active applications Q4 '26

S E Q U E N C I N G · T H E V O L I I I D E L T A

The five-framework convergence reorders the prior addendum: Bet #1 first (claim Process Power at Meta-B), Bet #2 second (collect equity at a power-holder), Bet #3 third (advisory practice compounding from both). Bets 4-5 fold into Bet #1 as modules; Bet #6 is the distribution layer; Bet #7 the long-arc fallback.

C R U X E S T H A T R E - R A N K E V E R Y T H I N G

- C1 · Anthropic ARR 24Bvs30B · Q3 '26
- C2 · MCP commons vs fork · EOY '26
- C3 · OSWorld 65% on frontier · Q3 '26
- C4 · EU AI Act Article 14 teeth · late '26
- C5 · Standalone memory absorbed vs niche · H2 '26